

Transaction

# Transaction

It is a set of operations used to perform a logical unit of work

# Types of operations

- Read
- Write
- Commit

# Example

$A=1000$   $B=2000$

$R(A)$

$A=A-500$

$W(A)$

$R(B)$

$B=B+500$

$W(B)$

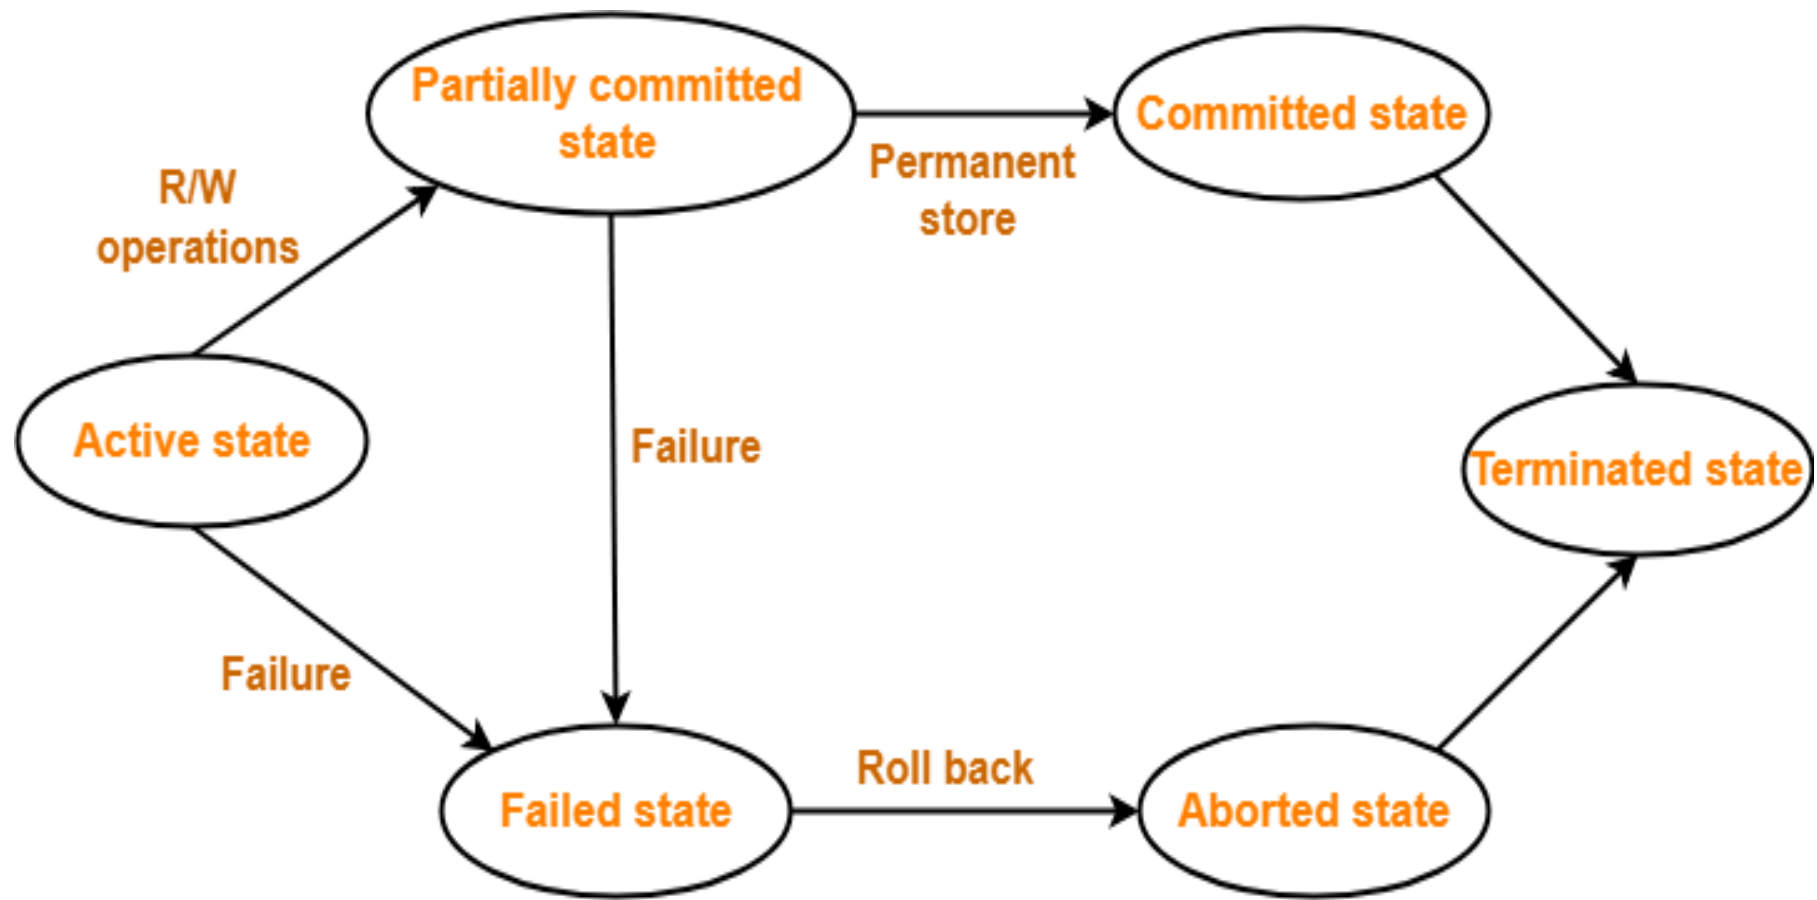
Commit----- $A=500$   $B=2500$

# ACID Properties

- Atomicity
- Consistency
- Isolation
- Durability

# Transaction States

- Active
- Partially Committed
- Committed
- Terminated
- Failed
- Abort



**Transaction States in DBMS**

# Schedule

It is a chronological execution sequence of multiple transactions.

OR

Multiple transactions executed in which order or sequence that is schedule.

Example:-  $T_1, T_2, T_3, \dots, T_n$ .



# Types

- Serial
- parallel

# Concurrency control protocol

Achieve serializability and recoverability-locking protocols.

## 1) Shared-Exclusive locking

- shared(S):-read only

- exclusive(X):-read and write both

eg. T1	T2
S(A)	X(A)
R(A)	R(A)
U(A)	W(A)
	U(A)

Compatibility table:-



YES	NO
NO	NO


# Advantages

- Simple protocol
- Simple programming
- Simple software

# Problems


May not sufficient to produce only serializable schedule.

Example,




# May not free from irrecoverability.

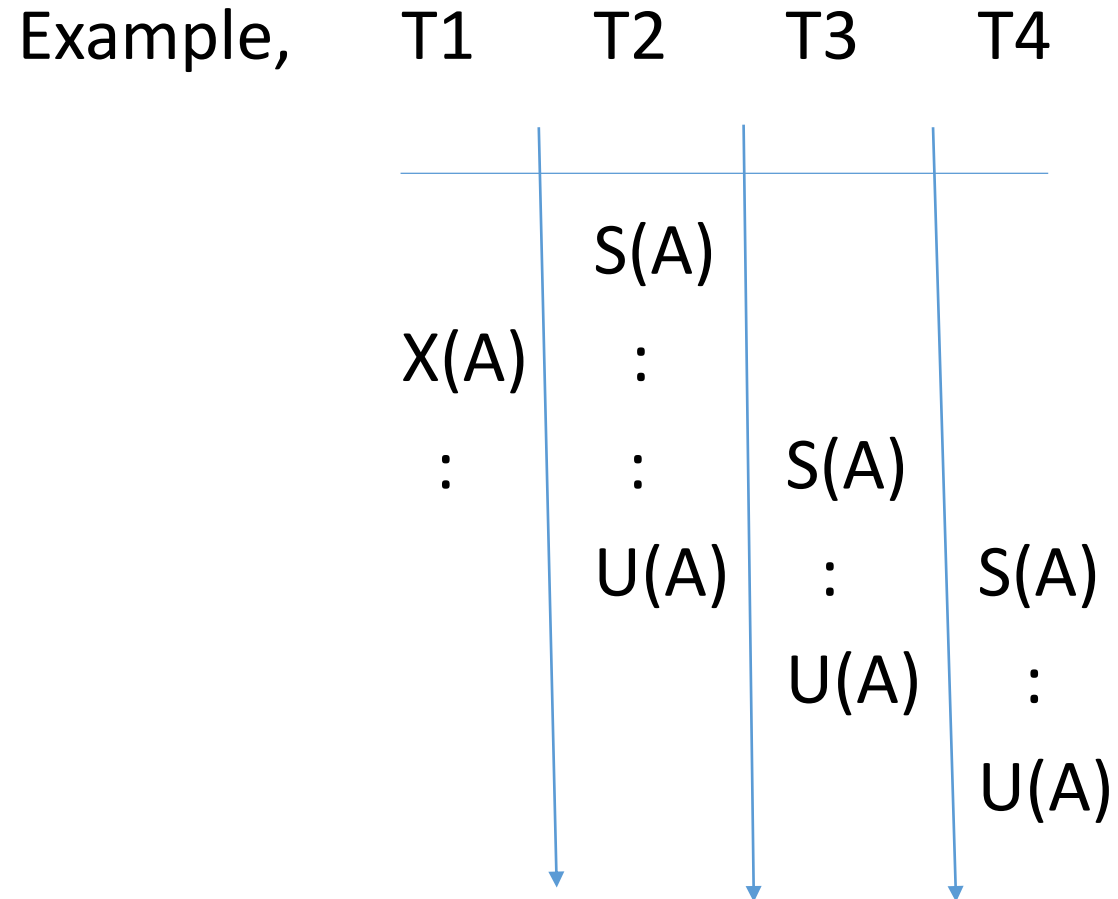
Example,



# May not free from deadlock



# May not free from starvation





## 2-phase locking(2PL)

- Growing:-locks are acquired and no locks are released.
- Shrinking:-locks are released and no locks are acquired.

# Example

T1

R(A)

W(A)

R(B)

R(A)

R(C)



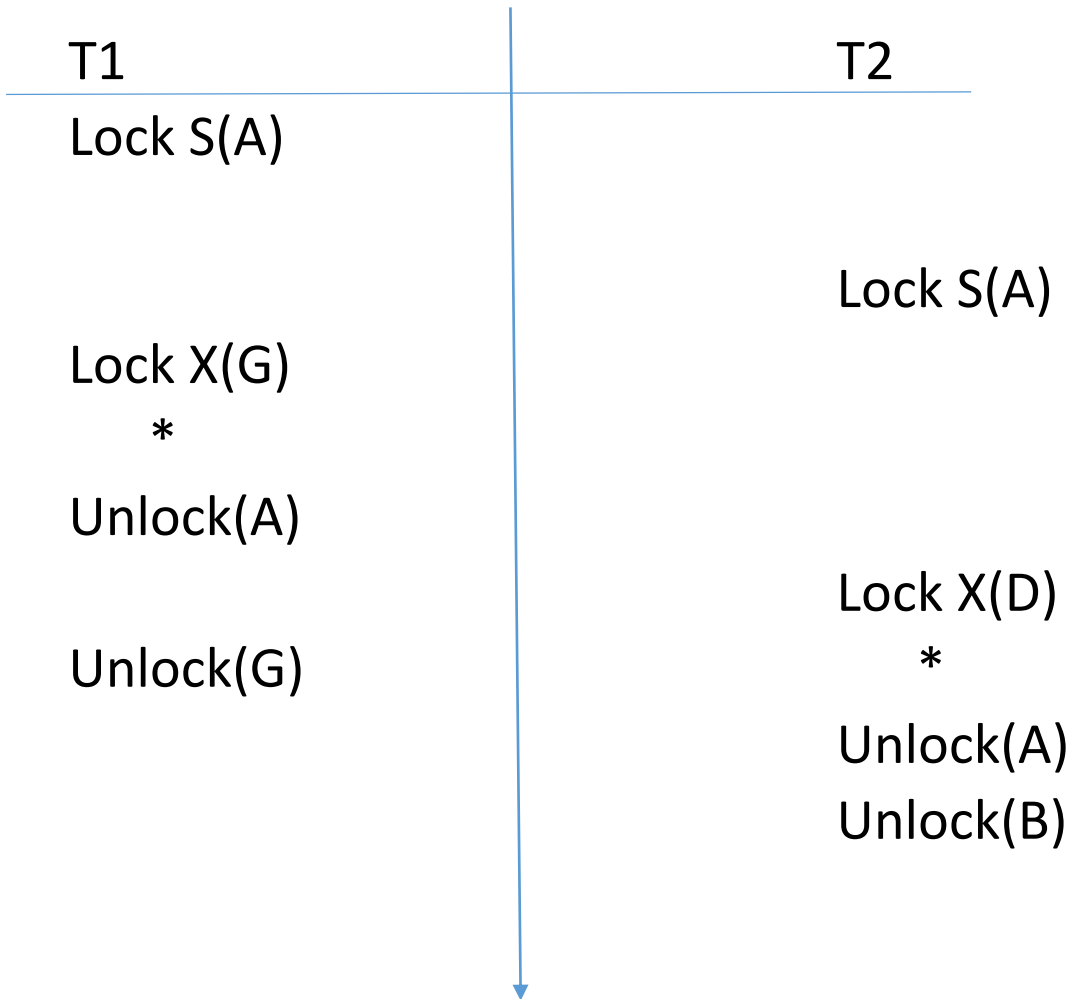
# Example

T1	T2
R(A) W(A)	
	R(A)
R(B) U(A) : : U(B)	

# Note

Any transaction that follows 2PL, it will always be serializable.

# Example




# Advantages:-

Always ensure serializability

# Drawbacks:-


- May not free from irrecoverability.

Example:-




- Not free from cascading rollback

Example:-






- Not free from deadlocks(Infinite wait)



- Not free from starvation (Finite wait)



Example:-




# Extensions of basic 2PL

- Strict 2PL:-It should satisfy the basic 2PL and all exclusive locks should hold until commit/Abort.
- Rigorous 2PL:-It should satisfy the basic 2PL and all shared, Exclusive locks should hold until commit/Abort. It is more restricted than strict 2PL.

# Example:-



# Example:-



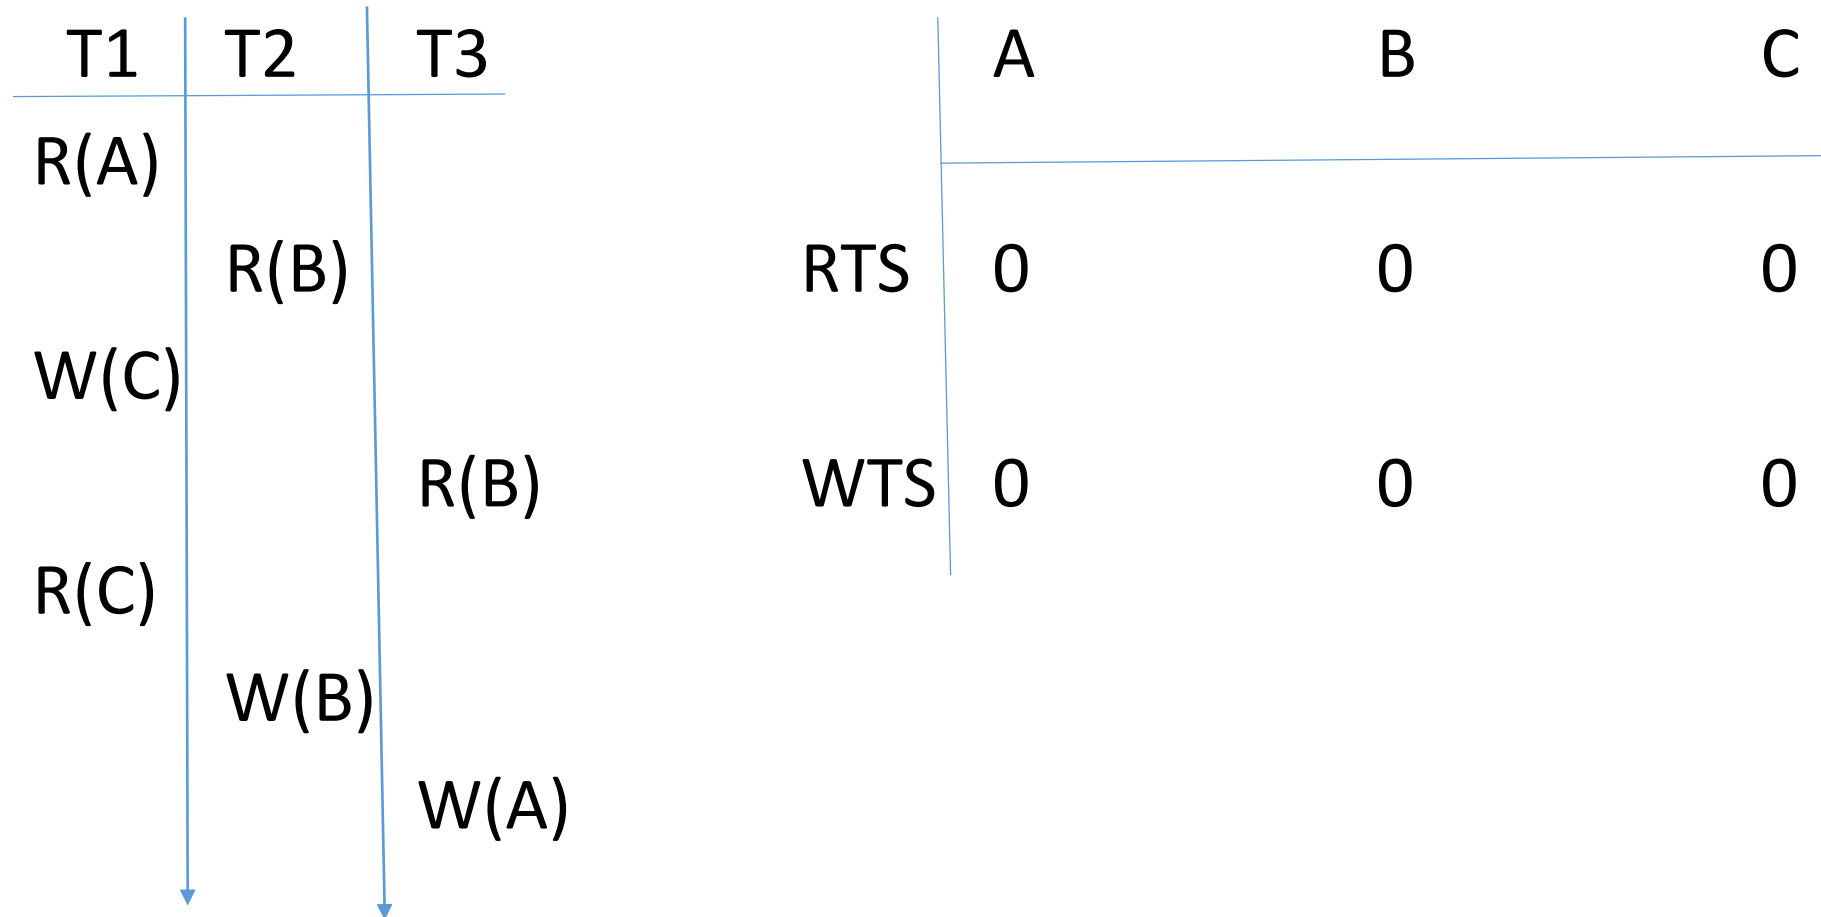
# Advantages:-

- Cascadeless
- Recoverable

# Drawbacks:-

- Deadlock
- Starvation

# Example-Timestamp ordering protocol





# Pessimistic concurrency control protocol

- Transaction delay
- Transaction rollback
- Unnecessary overhead

# Validation based or optimistic concurrency control protocol

- No checking is done while the transaction is executing.
- It is based on assumption that majority of database operations do not conflict.
- A transaction proceeds without restrictions until the transaction gets commit.
- It is then checked to see whether it has come into conflict with other transactions.
- When a conflict occurs, a transaction is aborted.

# Three phases

- Read phase
- Validation phase
- Write phase

# Read Phase

- At the start of this phase, transaction  $t_i$  is associated with a timestamp  $\text{start}(T_i)$ .
- $T_i$  reads the values of data items from the database and these values are then stored in the temporary local copies of the data items kept in the workspace of  $T_i$ .
- All modifications are performed on these temporary local copies of the data items without updating the actual data item of the database.

# Validation Phase

- At the start of this phase, transaction  $T_i$  is associated with a timestamp  $\text{validation}(T_i)$ .
- The system performs a validation test when  $T_i$  decides to commit.
- This validation is performed to find whether the modification made to be temporary local copies can be copied to the database without violating serializability.
- If  $T_i$  conflict with any other concurrently executing transaction,  $t_i$  is rolled back, its work place is cleared and  $T_i$  is restarted.

# Write Phase

- The system copies the modifications made by  $T_i$  in its workspace to the database only if it succeeds in the validation phase.
- At the end of this phase,  $T_i$  is associated with a timestamp  $\text{finish}(T_i)$ .

# Validation test for a transaction $T_j$

- For each transaction  $T_i$  with

$$Ts(T_i) < Ts(T_j)$$

Either one of the following condition holds:-

- $Finish(T_i) < start(T_j)$

Transaction  $T_j$  starts before  $T_i$  finishes  
( $\text{start}(T_j) < \text{finish}(T_i)$ ):-

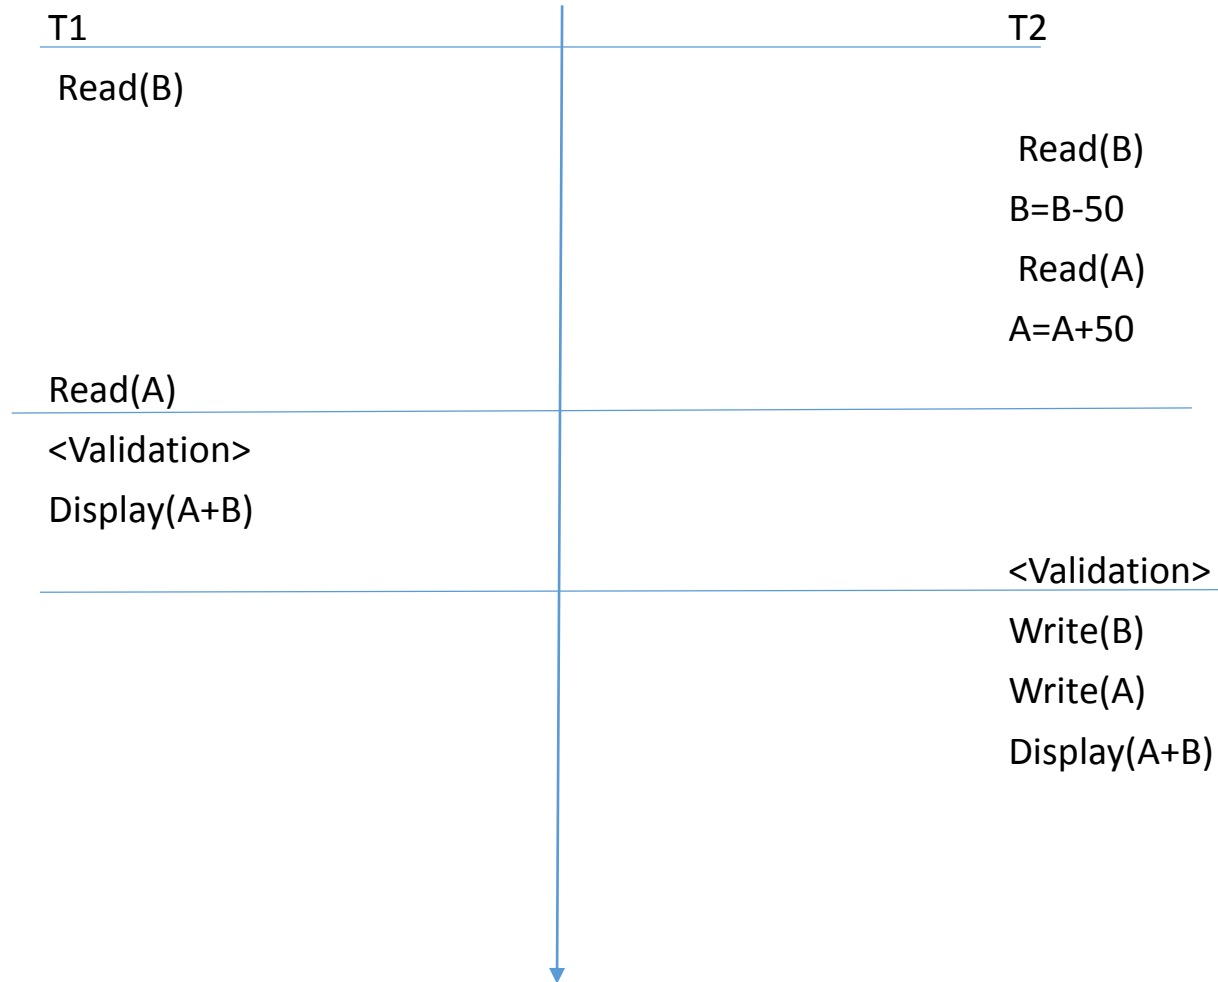
- $\text{Finish}(T_i) < \text{Validation}(T_j)$



# Transaction $T_j$ starts before $T_i$ finishes:-

- $\text{Validation}(T_i) < \text{Validation}(T_j)$

# Example



# Deadlock handling

- Prevention
- Detection

# Deadlock prevention

It ensures that deadlock never happens.

# Deadlock prevention techniques:-

- Wait-die technique
- Wound-wait technique

# Wait-die

(a) If a transaction  $T_i$  request for a resource that is locked by transaction  $T_j$ ,

if  $Ts(T_i) < Ts(T_j)$  then

$T_i$  is allowed to wait until resource is available for execution.

# Wait-die

(b) If  $T_i$  is older transaction and has hold some resources with it and  $T_j$  is waiting for it, then  $T_j$  is rolled back (dies).

if  $Ts(T_i) < Ts(T_j)$  then

$T_j$  is rolled back and restarted with very minute delay with same timestamp.

Hence, this technique allows the older transaction to wait but kills(dies) the younger transaction.

# Wound-wait

- (a) If transaction  $T_i$  request for a resource that is locked by transaction  $T_j$ ,
- if  $Ts(T_i) < Ts(T_j)$  then
- $T_j$  is rolled back (i.e wounded by  $T_i$ )



# Wound-wait

(b) If transaction  $T_j$ (younger) requesting a resource which is held by  $T_i$ (old),

if  $Ts(T_i) < Ts(T_j)$  then  $T_j$  waits

# Advantages of both techniques:-

- Starvation is avoided.
- A transaction with smallest timestamp is not rolled back.

## Disadvantage:-

- The request to acquire a lock on a data item held by another transaction does not necessarily involve a deadlock.
- Therefore, unnecessary rollbacks may occur in both wait-die and wound-wait techniques.

# Deadlock detection using wait-for graph

- It is suitable for smaller databases.
- To detect deadlock, the system maintains a wait-for graph, which consists of nodes and directed edges.
- Nodes of graph represent currently executing transactions.
- Directed edges exist from one node to another if a transaction is waiting for another transaction to release a lock.
- If this graph contains a cycle, it indicates the deadlock in the system.

# Failure classification

- Transaction failure
- System crash/Computer failure
- Disk failure
- Physical problem and environmental disasters

# Transaction failure

A transaction has to abort when it fails to execute or when it reaches a point from where it can't go any further is called transaction failure.

Reasons:-

- Logical error
- system error

# Storage structure

- Volatile
- Nonvolatile
- Stable

# Stable

a mythical form of storage that survives all failures.

approximated by maintaining multiple copies on distinct nonvolatile media.

- Maintain multiple copies of each block on separate disks
  - copies can be at remote sites to protect against disasters such as fire or flooding.
- Failure during data transfer can still result in inconsistent copies: Block transfer can result in
  - Successful completion
  - Partial failure: destination block has incorrect information
  - Total failure: destination block was never updated
- Protecting storage media from failure during data transfer (one solution):
  - Execute output operation as follows (assuming two copies of each block):
    1. Write the information onto the first physical block.
    2. When the first write successfully completes, write the same information onto the second physical block.
    3. The output is completed only after the second write successfully completes.



# stable

- Copies of a block may differ due to failure during output operation. To recover from failure:  
First find inconsistent blocks:
  1. *Expensive solution*: Compare the two copies of every disk block.
  2. *Better solution*:
    - Record in-progress disk writes on non-volatile storage (Non-volatile RAM or special area of disk).
    - Use this information during recovery to find blocks that may be inconsistent, and only compare copies of these.
    - Used in hardware RAID systems

# Recovery methods-Log based recovery

- Deferred database modification
- Immediate database modification

# Deferred database modification

T1	<T1,start>
R(A)	<T1,A,200>
A=A+100	<T1,B,400>
W(A)	<T1,commit>
R(B)	
B=B+200	
W(B)	
COMMIT	

# Example

<T1,start>

<T1,A,200>

<T1,B,400>

<T1,commit>

<T2,start>

<T2,C,500>

# Immediate database modification

T1	<T1,start>
R(A)	<T1,A,100,200>
A=A+100	<T1,B,200,400>
W(A)	<T1,commit>
R(B)	
B=B+200	
W(B)	
COMMIT	

# Example

<T1,start>

<T1,A,1000,2000>

<T1,B,5000,6000>

<T1,commit>

<T2,start>

<T2,C,700,800>