

# Natural Language Processing (NLP)

# Vector Space Model (VSM)

- Generalization of the Bag of Words model
- Each document from the corpus is represented as a multi- dimensional vector
  - Each unique term from the corpus represents one dimension of the vector space
  - *Term* can be a single word or a sequence of words (phrase)
  - The number of unique terms in the corpus determines the dimension of the vector space

# Vector Space Model (VSM)



- Vector elements are **weights** associated with individual terms;
- weights reflect the relevancy of the corresponding terms in the given corpus

# Vector Space Model (VSM)

- If a corpus consists of  $n$  terms  $(t_i, i=1, n)$ , document  $d$  from that corpus would be represented with the vector:  $d = \{w_1, w_2, \dots, w_n\}$ , where  $w_i$  are weights associated with terms  $t_i$

# Vector Space Model (VSM)

In VSM, corpus is represented in the form of ***Term Document Matrix (TDM)***, i.e., an  $m \times n$  matrix with following features:

- Rows ( $i=1,m$ ) represent terms from the corpus
- Columns ( $j=1,n$ ) represent documents from the corpus
- Cell  $ij$  stores the weight of the term  $i$  in the context of the document  $j$

# Vector Space Model (VSM)

Documents



Vector-space  
representation

However, complexity  
We will see how small  
Given a function based  
Using entropy of traffic  
We study the complexity  
of influencing elections  
through bribery: How  
computationally complex  
is it for an external actor  
to determine whether by  
a certain amount of  
bribing voters a specified  
candidate can be made  
the election's winner? We  
study this problem for  
election systems as varied  
as scoring ...

	D1	D2	D3	D4	D5
complexity	2		3	2	3
algorithm	3			4	4
entropy	1			2	
traffic		2	3		
network		1	4		

Term-document matrix

# Vector Space Model (VSM)

- Before creating the TDM matrix, documents from the corpus need to be *preprocessed*
- Rationale / objective: to reduce the set of words to those that are expected to be the *most relevant* for the given corpus

# Vector Space Model (VSM)

- Preprocessing (often) includes:
  - Normalizing the text
  - Removing terms with very small / high frequency in the given corpus
  - Removing the so-called stop-words
  - Reducing words to their root form through *stemming or lemmatization*



# Vector Space Model (VSM)

## Normalization Of Text

- **Objective:** transform various forms of the same term into a common, 'normalized' form
  - E.g.: Apple, apple, APPLE -> apple
  - Intelligent Systems, Intelligent systems, Intelligent-systems -> intelligent systems

# Vector Space Model (VSM)

## Normalization Of Text


### – How it is done:

- Using simple rules:
  - Remove all punctuation marks (dots, dashes, commas,...)
  - Transform all words to lower case
- Using a dictionary, such as **WordNet**, to replace synonyms with a common, often more general, concept
  - E.g., “automobile, car” -> vehicle

# Vector Space Model (VSM)

## Removing High And Low Frequency Terms

- Empirical observations (in numerous corpora):
  - Many low frequency words
  - Only a few words with high frequency
- Formalized in the ***Zipf's rule***:
  - the frequency of a word in a given corpus is inversely proportional to its rank in the frequency table (for that corpus)




# Vector Space Model (VSM)

## Implications Of The Zipf's Rule

- high frequency but are semantically almost useless
  - Examples: the, a, an, we, do, to
- semantically rich, but are of very low frequency
  - Example: dextrosinistral
- The rest of the words are those that represent the corpus the **best** and thus should be **included** in the VSM model

Note : dextrosinistral=expert in physical movement with hands



# Vector Space Model (VSM)

- Stop-words are those words that (on their own) do not bear any information / meaning
- It is estimated that they represent 20-30% of words in any corpus
- There is no unique stop-words list
  - Frequently used lists are available at:  
<http://www.ranks.nl/stopwords>
- Potential problems with stop-words removal:
  - the loss of original meaning and structure of text
  - examples: “this is not a good option” -> “option”
  - “to be or not to be” -> null

# **Lemmatization And Stemming ?**

**Already discussed in tutorial**



# Vector Space Model (VSM)

## Computing Terms' Weights

- Simple and frequently used approaches include:
  - Binary weights
  - Term Frequency (TF)
  - Inverse Document Frequency (IDF)
  - TF-IDF

# Vector Space Model (VSM)

## Binary Weights

- Weights take the value of 0 or 1, to reflect the presence (1) or absence (0) of the term in a particular document

# Vector Space Model (VSM)

## Binary Weights

	text	information	identify	mining	mined	is	useful	to	from	apple	delicious
Doc1	1	1	1	1	0	1	1	1	0	0	0
Doc2	1	1	0	0	1	1	1	0	1	0	0
Doc3	0	0	0	0	0	1	0	0	0	1	1

### Example:

Doc1: Text mining is to identify useful information.

Doc2: Useful information is mined from text.

Doc3: Apple is delicious.

# Vector Space Model (VSM)

## Term Frequency

TF represents the frequency of the term in a specific document

- The underlying assumption: the higher the term frequency in a document, the more important it is for that document
- $TF(t) = c(t,d)$
- $c(t,d)$  – the number of occurrences of the term  $t$  in the document  $d$

# Vector Space Model (VSM)

## Inverse Document Frequency (IDF)

- Assign higher weights to unusual terms, i.e., to terms that are not so common in the corpus
- IDF is computed at the **corpus level**, and thus describes corpus as a whole, not individual documents
- It is computed in the following way:

$$\text{IDF}(t) = 1 + \log(N/df(t))$$

Where,  $N$  = number of documents in the corpus

$df(t)$  = number of documents with the term  $t$

# Vector Space Model (VSM)

## TF-IDF

- The underlying idea: value those terms that are not so common in the corpus (relatively high IDF), but still have same reasonable level of frequency (relatively high TF)
- The **most frequently** used metric for computing term weights in a VSM
- One popular ‘instantiation’ of this formula:  
$$\text{TF-IDF}(t) = \text{tf}(t) * \log(N/\text{df}(t))$$

# Vector Space Model (VSM)


## Estimating **Similarity** Of Documents

- Key question: which metric to use for estimating the similarity of documents (i.e., vectors that represent documents)?
- The most well known and widely used metric is ***Cosine similarity*** “It is dot product of two vectors divided by the product of the two vector’s lengths (or magnitudes).”
- Cosine Similarity

$$\cos(d_i, d_j) = V_i \cdot V_j / (||V_i|| \cdot ||V_j||)$$


$V_i$  and  $V_j$  are vectors representing documents  $d_i$  and  $d_j$

# Vector Space Model (VSM)







# Vector Space Model (VSM)



Similar scores  
Score Vectors in same direction  
Angle between them is near 0 deg.  
Cosine of angle is near 1 i.e. 100%



Unrelated scores  
Score Vectors are nearly orthogonal  
Angle between them is near 90 deg.  
Cosine of angle is near 0 i.e. 0%



Opposite scores  
Score Vectors in opposite direction  
Angle between them is near 180 deg.  
Cosine of angle is near -1 i.e. -100%